

0.1 A simple GUI in R

We begin with an example showing how one can use R's standard graphics device as a canvas for a "game" of tic-tac-toe against the computer (Figure 0.1). Although this example has nothing to do with statistics, it illustrates, in a familiar way, some of the issues involved in developing GUIs in R.

Generally, GUIs provide the means for viewing and controlling some underlying data structure. In this example, the data simply consists of information holding the state of the game, defined here in a global variable `board`.

```
board <- matrix(rep(0,9), nrow=3)
```

A GUI contains one or more views, each of which displays the data in a particular manner. In our case, the view is the game board that we display through an R graphics device. The `layoutBoard` function creates a canvas for this view:

```
layoutBoard <- function() {  
  plot.new()  
  plot.window(xlim=c(1,4), ylim=c(1,4))  
  abline(v=2:3); abline(h=2:3)  
  mtext("Tic Tac Toe. Click a square:")  
}
```

This example uses a single view; more complex GUIs will contain multiple coordinated, interactive views. The layout of the GUI should help the user navigate the interface and is an important factor in usability. Here we benefit from the universal familiarity with the board game.

The user typically sends input to a GUI through a mouse or keyboard. The underlying toolkit allows the programmer to assign functions to be called when some specific event occurs, such as user interaction. Typically, the toolkit *signals* that some action has occurred, and then invokes *callbacks* or *event handlers* that have been assigned by the programmer. Each toolkit has a different implementation. For our game, we will use the `locator` function built into the base R graphics system:

```
doPlay <- function() {  
  iloc <- locator(n=1, type="n")  
  clickHandler(iloc)  
}
```

The `locator` function responds to mouse clicks. One specifies how many mouse clicks to gather and the *control* of the program is suspended until the user clicks the sufficient number of times (or somehow interrupts the loop). Such a GUI that enters a mode in which the flow of a program is blocked and waiting on user input is known as a *modal* GUI. This design is

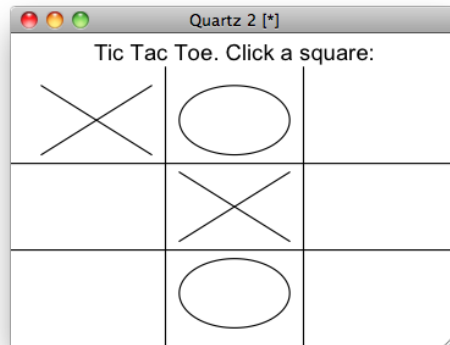


Figure 0.1: Using a graphics device for a game of tic-tac-toe

common for simple dialogs that require immediate user attention, although in general a GUI will listen asynchronously for user input.

In the above function `doPlay`, `clickHandler` is an *event handler*. Its job is to process the output of the `locator` function, checking first if the user terminated `locator` using the keyboard. If not it proceeds to draw the move, and then, if necessary, the computer's move. Afterwards, play is repeated until there is a winner or a "cat's" game.

```
clickHandler <- function(iloc) {  
  if(is.null(iloc))  
    stop("Game terminated early")  
  move <- floor(unlist(iloc))  
  drawMove(move, "x")  
  board[3*(move[2]-1) + move[1]] <- 1  
  if(!isFinished())  
    doComputerMove()  
  if(!isFinished())  
    doPlay()  
}
```

The use of `<-` in the handler illustrates a typical issue in GUI design in R. User input changes the state of the application through callback functions. These callbacks need to modify variables in some shared scope, which may be application-wide or specific to a component. The lexical scoping rules of R, i.e., nesting of closures, has proven to be a useful strategy for managing GUI state. In the above case, we simply modify the global environment, which encloses `clickHandler`. When this is inconvenient, direct manipulation of environment objects is sometimes a feasible

option. If the scale of the GUI demands more formal mechanisms, we recommend the reference class framework from the `methods` package.

Validation of user input is an important task for a GUI. In the above, the `clickHandler` function checks to see if the user terminated the game early and issues a message.

At this point, we have a data model, a view of the model and the logic that connects the two, but we still need to implement some of the functions to tie it together.

This function draws either an “x” or an “o”. It does so using the `lines` function. The `z` argument contains the coordinates of the square to draw.

```
drawMove <- function(z,type="x") {
  i <- max(1,min(3,z[1])); j <- max(1,min(3,z[2]))
  if(type == "x") {
    lines(i + c(.1,.9),j + c(.1,.9))
    lines(i + c(.1,.9),j + c(.9,.1))
  } else {
    theta <- seq(0,2*pi,length=100)
    lines(i + 1/2 + .4*cos(theta), j + 1/2 + .4*sin(theta))
  }
}
```

One could use text to place a text “x” or “o”, but this may not scale well if the GUI is resized. Most GUI layouts allow for dynamic resizing. This is necessary to handle the variety of data a GUI will display. Even the labels, which one generally considers static, will display different text depending on the language (as long as translations are available).

This function is used to test if a game is finished:

```
isFinished <- function() {
  (any(abs(rowSums(board)) == 3) ||
   any(abs(colSums(board)) == 3) ||
   abs(sum(diag(board))) == 3 ||
   abs(sum(diag(apply(board, 2, rev)))) == 3)
}
```

The matrix `m` allows us to easily check all the possible ways to get three in a row.

This function picks a move for the computer:

```
doComputerMove <- function() {
  newMove <- sample(which(board == 0),1) # random !
  board[newMove] <- -1
  z <- c((newMove-1) %% 3, (newMove-1) %/% 3) + 1
  drawMove(z,"o")
}
```

The move is converted into coordinates using `%` to get the remainder and `%%` to get the quotient when dividing an integer by an integer. This function just chooses at random from the left over positions; we leave implementing a better strategy to the interested reader.

Finally, we implement the main entry point for our GUI:

```
playGame <- function() {  
  board <- matrix(rep(0,9), nrow=3)  
  layoutBoard()  
  doPlay()  
  mtext("All done\n",1)  
}
```

When the game is launched, we first lay out the board and then call `doPlay`. When `doPlay` returns, a message is written on the screen.

This example adheres to the model-view-controller design pattern that is implemented by virtually every complex GUI. We will encounter this pattern throughout this book, although it is not always explicit.

For many GUIs there is a necessary trade-off between usability and complexity. As with any software, there is always the temptation to continually add features without regard for the long term cost. In this case, there are many obvious improvements: implementing a better artificial intelligence, drawing a line connecting three in a row when there is a win, indicating who won, etc. Adding a feature increases the functionality, at the cost of increased complexity and burden on the user.

0.2 GUI Design Principles

The most prevalent pattern of user interface design is denoted WIMP, which stands for Window, Icon, Menu and Pointer. The WIMP approach was developed at Xerox PARC in the 1970's and later popularized by the Apple Macintosh in 1984. This is particularly evident in the separation of the window from the menubar on the Mac desktop. Other graphical operating systems, such as Microsoft Windows, later adapted the WIMP paradigm, and libraries of reusable GUI components emerged to support development of applications in such environments. Thus, GUI development in R adheres to the WIMP approach.

The primary WIMP component from our perspective is the window. A typical interface design consists of a top-level window referred to as the *document window* that shows the current state of a “document,” whatever that is taken to be. In R it could be a data frame, a command line, a function editor, a graphic or an arbitrarily complex form containing an assortment of such elements.

Abstractly, WIMP is a command language, where the user executes commands, often called actions, on a document by interacting with graph-

ical controls. Every control in a window belongs to some abstract menu. Two common ways of organizing controls into menus are the menubar and toolbar.

The parameters of an action call, if any, are controlled in sub-windows. These sub-windows are termed *application windows* by Apple^[1], but we prefer the term *dialogs*, or *dialog boxes*. These terms may also refer to smaller sub-windows that are used for alerts or confirmation. The program often needs to wait for user input before continuing with an action, in which case the window is modal. We refer to these as *modal dialog boxes*.

Each window or dialog typically consists of numerous controls laid out in some manner to facilitate the user interaction. Each window and control is a type of *widget*, the basic element of a GUI. Every GUI is constituted by its widgets. Not all widgets are directly visible by the user; for example, many GUI frameworks employ invisible widgets to lay out the other widgets in a window.

There is a wide variety of available widget types, and widgets may be combined in an infinite number of ways. Thus, there are often numerous means to achieve the same goals. For example, Figures 0.2 and 0.3 show three dialogs, representing typical dialogs from the three main operating systems, that perform the same task – collect arguments from the user to customize the printing of a document. Although all were designed to do the same thing, there are many differences in implementation.

In some cases, typical usage suggests one control over another. The choice of printer for each is specified through a combo box. However, for other choices a variety of widgets are employed. For example, the control to indicate the number of copies for the Mac is a simple text entry window, whereas for the KDE and Windows dialog it is a spin button. The latter provides a bit more functionality, for a bit more complexity. The KDE and Mac dialogs have icons to compactly represent actions, whereas the the Windows example has none. The landscape icon for the Mac is very clear and provides this feature without having to use a sub dialog.

How the interfaces are laid out also varies. All panels are read top to bottom, although the Mac interface also has a very nice preview feature on the left side. The KDE dialog uses frames to separate out the printer arguments from the arguments that specify how the print job is to proceed. The Mac uses a vertical arrangement to guide the user through this. For the Mac, horizontal separators are used instead of frames to break up the areas, although a frame is used towards the bottom. Apple uses a center balance for its controls. They are not left justified as are the KDE and Windows dialogs. Apple has strict user-interface guidelines and this center balance is a design decision.

[1] Apple Inc. <http://developer.apple.com/>.

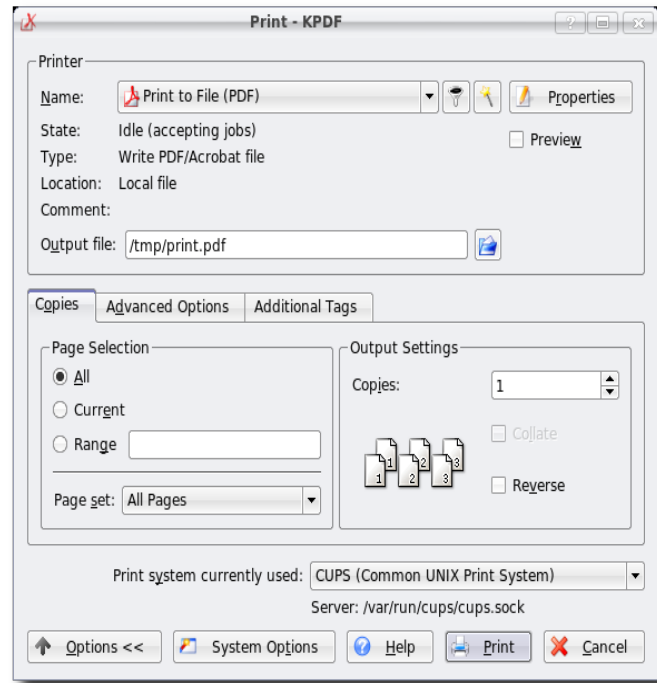
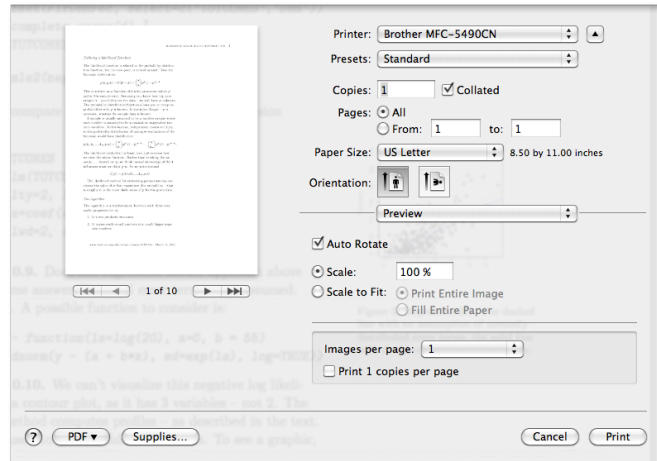


Figure 0.2: Two print dialogs. One from Mac OS X 10.6 and one from KDE 3.5.

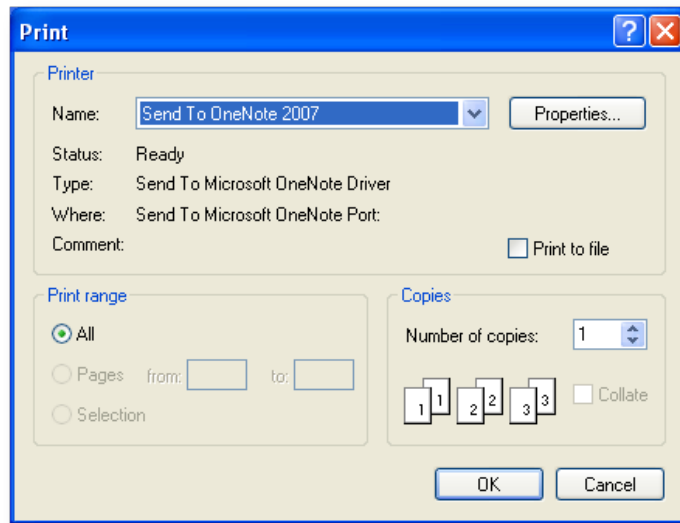


Figure 0.3: R's print dialog under Windows XP using XP's native dialog.

The layout also determines how many features and choices are visible to the user at a given time. For example, the Mac GUI uses “disclosure buttons” to allow access to printer properties and the PDF settings, whereas KDE uses a notebook container to show only a subset of the options at once.

The Mac GUI provides a very nice preview of the current document indicating, to the user clearly what is to be printed and how much. Adjusting GUIs to the possible state is an important user interface property. GUI areas that are not currently sensitive to user input are grayed out. For example, the “collate” feature of the GUI only makes sense when multiple copies are selected, so the designers have it grayed out until then. A common element of GUI design is to only enable controls when their associated action is possible, given the state of the application.

The Mac GUI has the number of pages in focus, whereas Windows places the printer in focus. Focus allows the user to interact with the GUI without the mouse. Typically the tab key is used to step through the controls. GUI's often have shortcuts that allow power users to initiate actions or shift the focus directly to a specific widget through the keyboard. Most dialogs also have a default button, which will initiate the dialog action when the return key is pressed. The KDE dialog, for example, indicates that the “print” button is the default button through special shading.

Each dialog presents the user with a range of buttons to initiate or cancel the printing. The Windows ones are set on the right and consist of

Table 0.1: Table of possible selection widgets by data type, size and selection mode (single or multiple)

Type of data	Single	Multiple
Boolean	Checkbox, toggle button	-
Small list	radio button group	checkbox group
	combo box	list box
	list box	
Moderate list	combo box	list box
	list box	
Large list	list box, auto complete	list box
Sequential	slider	
	spin button	
Tabular	table	table
Hierarchical	tree	tree

the standard “OK” and “Cancel” buttons. The Mac interface uses a spring to push some buttons to the left, and some to the right to keep separate their level of importance. The KDE buttons do so as well, although one cannot tell from the image. The use of conventional icons on the buttons also help guide the user.

0.3 Controls

This section provides an overview of many common controls, i.e., widgets that either accept input, display data or provide visual guides to help the user navigate the interface. If the reader is already familiar with the conventional types of widgets and how they are arranged on the screen, this section and the next should be considered optional.

Choice of control

A GUI is comprised of one or more widgets. The appropriate choice depends on a balance of considerations. For example, many widgets offer the user a selection from one or more possible choices. An appropriate choice depends on the type and size of the information being displayed, the constraints on the user input, and on the space available in the layout. As an example, Table 0.3 suggests different types of widgets used for this purpose depending on the type and size of data and the number of items to select.

Figure 0.4 shows several such controls in a single dialog. A checkbox enables an intercept, a radio group selects either full factorial or a custom

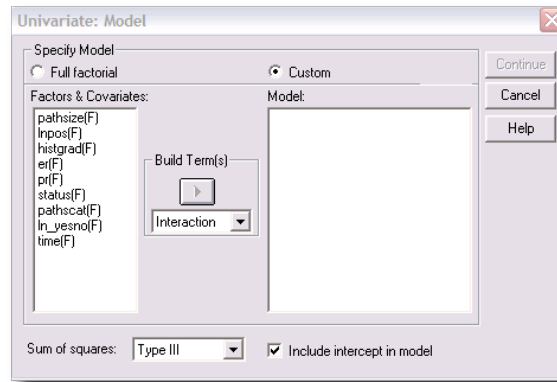


Figure 0.4: A dialog box from SPSS version 11 for specifying terms for a linear model. The graphic shows a dialog that allows the user to specify individual terms in the model using several types of widgets for selection of values, such as a radio button group, a checkbox, combo boxes, and list boxes.

model, a combo box selects the “sum of squares” type, and a list box allows for multiple selection from the available variables in the data set.

For many R object types there are natural choices of widget. For example, values from a sequence map naturally to a slider or spin button; a data frame maps naturally to a table widget; or a list with similar structure can map naturally to a tree widget. However, certain R types have less common metaphors. For instance, a formula object can be fairly complex. Figure 0.4 shows an SPSS dialog to build terms in a model. R power users may be much faster specifying the formula through a text entry box, but beginning R users coming to grips with the command line and the concept of a formula may benefit from the assistance of a well designed GUI. One might desire an interface that balances the needs of both types of user, or the SPSS interface may be appropriate. Knowing the potential user base is important.

Presenting options

The widgets that receive user input need to translate that input into a command that modifies the state of the application. Commands, like R functions, often have parameters, or options. For many options, there is a discrete set of possible choices, and the user needs to select one of them. Examples include selecting a data frame from a list of data frames, selecting a variable in a data frame, selecting certain cases in a data frame, selecting a logical value for a function argument, selecting a numeric value for a

confidence level or selecting a string to specify an alternative hypothesis. Clearly there can be no one-size-fits-all widget to handle the selection of a value.

Checkboxes

A *checkbox* specifies a value for a logical (boolean) option. Checkboxes have labels to indicate which variable is being selected. Combining multiple checkboxes into a group allows for the selection of one or more values at a time.

Radio buttons

A *radio button group* selects exactly one value from a vector of possible values. The analogy dates back to old car radios where there were a handful of buttons to select a preset channel. When a new button was pushed in, the previously pressed button popped up. Radio button groups are useful, provided there are not too many values to choose from, as all the values are shown. These values can be arranged in a row, a column or both rows and columns to better fill the available space. Figure 0.5 uses radio button groups for choosing the distribution, kernel and sample size for the density plot.

Combo boxes

A *combo box* is similar to a radio button group, in that it is used to select one value from several. However, a combo box only displays the value currently selected, which reduces visual complexity and saves space, at the cost of an extra click to show the choices. Toolkits often combine a combo box with a text entry area for specifying an arbitrary value, possibly one that is not represented in the set of choices. A combo box is generally desirable over radio buttons when there are more than four or five choices. However, the combo box also has its limits. For example, some web forms require choosing a country from a list of hundreds. In such cases, features like incremental type ahead search are useful.

List boxes

A *list box* displays a list of possible choices in a column. While the radio button group and combo box select only a single value, a list box supports multiple selection. Another difference is that the number of displayed choices depends dynamically on the available space. If a list box contains too many items to display them simultaneously, a scrollbar is typically provided for adjusting the visible range. Unlike the combo box, the choices are immediately visible to the user. Figure 0.5 shows a list box created by

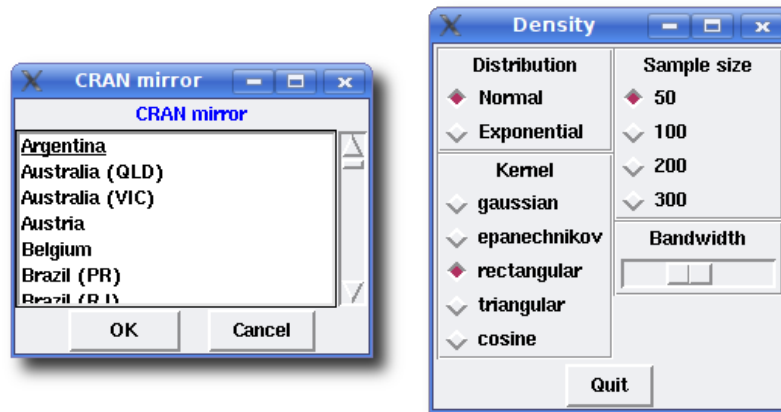


Figure 0.5: Two applications of the `tcltk` package. The left graphic is produced by `chooseCRANmirror` and uses a list box to allow selection from a long list of possibilities. The right graphic is the `tkdensity` demo from the `tcltk` package. It uses radio buttons and a slider to select the parameter values for a density plot.

the R function `chooseCRANmirror`. There are too many mirrors to fit on the screen, but a combo box would not take advantage of the available space. The list box is a reasonable compromise.

Sliders and spin buttons

A *slider* is a widget that selects a value from a sequence of possible values typically through the manipulation of a knob that moves or “slides” along a line that represents the range of possible values. Some toolkits generalize beyond a numeric sequence. The slider is a good choice for offering the user a selection of ordinal or numerical parameter values. For example, the letters of the alphabet could be a sequence. The `tkdensity` demo of the `tcltk` package (Figure 0.5) uses a slider to dynamically adjust the bandwidth of a density estimate.

A *spin button* plays a similar role to the slider, in that it selects a value within a set of bounds. Typically, this widget is drawn with a text box displaying the current value and two arrows to increment or decrement the selection. Usually, the text box can be edited directly. A spin button has the advantage of using less screen space, and directly entering a specific value, if known, is easier than selecting it with a slider. One disadvantage is that the position of the selected value within the range is not as obvious compared to the slider. As a compromise, combining a text box with a

slider is possible and often effective. A spin button is used in the KDE print dialog of Figure 0.2 to adjust the number of copies.

Initiating an action

After the user has specified the parameters of an action, typically by interacting with the selection widgets presented above, it comes time to execute the action. Widgets that execute actions include the familiar buttons, which are often organized into menubars and toolbars.

Buttons

A *button* issues commands when invoked, usually via a mouse click. In Figure 0.2, the “Properties” button, when clicked, opens a dialog for setting printer properties. The button with the wizard icon also opens a dialog. As buttons execute an action, they are often labeled with a verb.^[1] In Figure 0.4 we see how SPSS uses buttons in its dialogs: buttons which are not valid in the current state are disabled; buttons which are designed to open subsequent dialogs have trailing dots; and the standard actions of resetting the data, canceling the dialog or requesting help are given their own buttons on the right edge of the dialog box.

To speed the user through a dialog, a button may be singled out as the default button, so its action will be called if the user presses the return key. Actions may be given shortcut bindings, and their button proxies typically reflect the proper key combination to invoke the action. The KDE print dialog in Figure 0.2 has these bindings indicated through the underlined letter on the button labels.

Icons

In the WIMP paradigm, an *icon* is a pictorial representation of a resource, such as a document or program, or, more generally, a concept, such as a type of file. An application GUI typically adopts the more general definition, where an icon is used to complement or replace a text label on a button or other control. A button represents an action, so an icon on a button should visually depict an action.

Menu Bars

Menus play a central role in the WIMP desktop. The *menu bar* contains items for many of the actions supported by the application. By convention, menubars are associated with a top-level window. This is enforced by some toolkits and operating systems, but not all. In Mac OS X, the menubar appears on the top line of the display, but other platforms place

the menubar at the top of the top-level window. In a statistics application, the “document” may be the active data frame, a report, or a graphic.

The styles used for menubars are fairly standardized, as this allows new users to quickly orient themselves within a GUI. The visible menu names are often in the order File, Edit, View, Tools, then application specific menus, and finally a Help menu. Each visible menu item when clicked opens a menu of possible actions. The text for these actions conventionally use a ... to indicate that a subsequent dialog will open so that more information can be gathered to complete the action. The text may also indicate a keyboard accelerator, such as Find Next F3 indicating that both “N” as a keyboard accelerator and F3 as a shortcut will initiate this same action. (Shortcuts are not translated, but keyboard accelerators must be. As such, their use is less so. In particular, keyboard accelerators are not supported in Mac OS X menus.)

Not all actions will be applicable at any given time. It is recommended that rather than deleting these menu items, they be disabled, or grayed out, instead.

Menus may come to contain many items. To help the user navigate, menu items are usually grouped with either horizontal separators or hierarchical submenus.

The use of menus has evolved to also allow the user to view and control properties of the application state. There may be checkboxes drawn next to the menu item or some icon indicating the current state.

Another use of menus is to bind contextual menus (popup menus) to certain mouse clicks on GUI elements. Typically right mouse clicks will pop up a menu that lists often-used commands that are appropriate for that widget and the current state of the GUI. In Mac OS X one-button users, these menus are bound to a control-click.

Toolbars

Toolbars are used to give immediate access to the frequently used actions defined in the menubar. Toolbars typically have icons representing the action and perhaps accompanying text. They traditionally appear on the top of a window, but sometimes are used along the edges.

Action Objects

When clicking on a button, the user expects some “action” to occur. For example, some save dialog is summoned, or some page is printed. GUI toolkits commonly represent such actions as formal, invisible objects that are proxied by widgets, usually buttons, on the screen. Often, all of the primary commands supported by an application have a corresponding ac-

tion object, and the buttons associated with those actions are organized into menubars and toolbars.

An action object is essentially a data model, with each proxy widget acting as a view. Common components of an action include a textual label, an icon, perhaps a shortcut, and a handler to call when the action is selected.

Modal dialogs

A *modal dialog box* is a dialog box that keeps the focus until the user takes an action to dismiss the box. It prompts a user for immediate input, for example asking for confirmation when overwriting a file. Modal dialog boxes can be disruptive to the flow of interaction, so are used sparingly. As the control flow is blocked until the window is dismissed, functions that display modal dialogs can return a value when an event occurs, rather than have a handler respond to asynchronous input. The `file.choose` function, mentioned below, is a good example. When called during an interactive R session, the user is unable to interact with the command line until a file has been specified or the dialog dismissed.

Message dialogs

A *message dialog* is a high-level dialog widget for communicating a message to the user. By convention, there is a small rectangular box that appears in the middle of the screen with an icon on the left and a message on the right. At the bottom is a button to dismiss the dialog, often labeled “OK.” Additional buttons/responses are possible. The *confirmation dialog* variant would add a “Cancel” button which invalidates the proposed action.

File choosers

A file chooser allows for the selection of files and directories. They are familiar to any user of a GUI. A typical R installation has the functions `file.choose` and `tkchooseDirectory` (in the `tcltk` package) to select files and directories.

Other common choosers are color choosers and font choosers.

Displaying data

Table and tree widgets support the display and manipulation of tabular and hierarchical data, respectively. More arbitrary data visualization, such as statistical plots, can be drawn within a GUI window. All the toolkits we discuss have some means to embed R’s graphics.

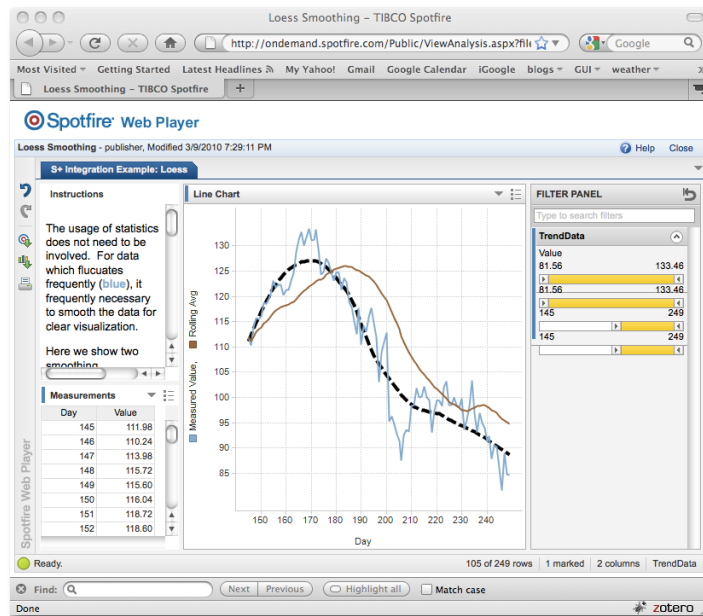


Figure 0.6: A screen shot from Tibco's Spotfire web player illustrating a table widget (lower left), displaying the cases that are summarized in the graphic. The right bar filters the cases in the table.

Tabular display

A *table widget* shows tabular data, such as a data frame, where each column has a specific data type and cell rendering strategy. Table widgets handle the display, sorting and selection of records from a dataset. Depending on the configuration of the widget, cells may be editable. Figure 0.6 shows a table widget in a Spotfire web player demonstration.

Tree widgets

So far, we have seen how list boxes display homogeneous vectors of data, and how table widgets display tabular data, like that in a data frame. Other widgets support the display of more complex data structures. If the data has a hierarchical structure, then a *tree widget* may be appropriate for its display. Examples of hierarchical data in R are directory structures, the components of a list, or class hierarchies. The object browser in JGR uses a tree widget to show the components of the objects in a user session (Figure 0.7). The root node of the tree is the "data" folder, and each data object in the global workspace is treated as an offspring of this root node. For the data frame *iraq*, its variables are considered as offspring of

the data frame. In this case these variables have no further offspring, as indicated by the “page” icon.

Displaying and editing text

The letter P in WIMP stands for “pointer,” so it is unsurprising that WIMP GUIs are designed around the pointing device. The keyboard is generally relegated to a secondary role, in part because it is difficult to both type and move the mouse at the same time. For statistical GUIs, especially when integrating with the command-line interface of R, the flexibility afforded by arbitrary text entry is essential for any moderately complex GUI. Toolkits generally provide separate widgets for text entry depending on whether the editor supports a single line or multiple lines.

Single line text

A text entry widget for editing a single line of text is found in the KDE print dialog (Figure 0.2). It specifies the page range. Specifying a complex page range, which might include gaps, would require a complex point-and-click interface. In order to avoid complicating the GUI for a feature that is rarely useful, a simple language has been developed for specifying page ranges. There is overhead involved in the parsing and validation of such a language, but it is still preferable to the alternative.

Text edit boxes

Figure 0.8 shows three multi-line text entries in an Rcmdr window. It provides an R console and status message area. The “Output Window” demonstrates the utility of formatting attributes. In this case, attributes specify the color of the commands, so that the input can be distinguished from the output.

Guides and feedback

Some widgets display information but do not respond to user input. Their main purpose is to guide and the user through the GUI and to display feedback and status messages. Communicating application status, such as during long-running calculations or when errors occur, is an often overlooked but critically important feature of any effective GUI.

Labels

A label is a widget for placing text into a GUI that is typically not intended for editing, or even selecting with a mouse. The main role of a label is to

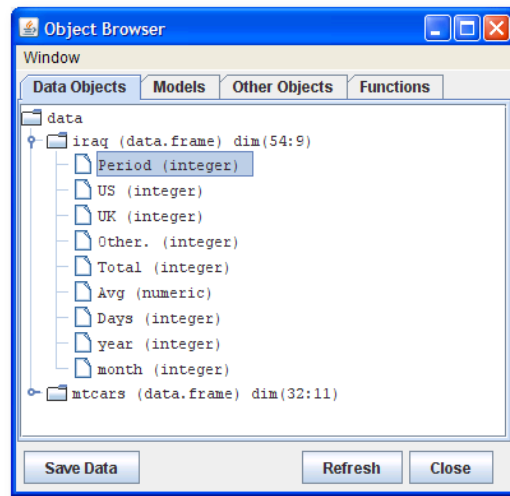


Figure 0.7: The object browser in the JGR GUI using a tree widget to display the possibly hierarchical nature of R objects.

describe another component of the GUI. Most toolkits support rich text in labels. Figure 0.8 shows labels marked in red and blue in `tc1tk`.

Statusbars

A statusbar displays general status messages, as well as feedback on actions initiated by the user, such as progress or errors. Messages replace the previous message and may disappear after a certain period of time. In the traditional document-oriented GUI, statusbars are placed at the bottom.

Related to status bars are info bars or alert boxes, that allow a programmer to display a transient message dialog that emerges from one either the top or bottom the application window. An example is the Firefox dialog that asks whether Firefox should remember a password entered on the previous page. It appears just below the toolbar and disappears automatically as the user continues to browse.

Tooltips

A tooltip is a small window that is displayed when a user hovers their mouse over a tooltip-enabled widget. These are an embellishment for providing extra information about a particular piece of content displayed by a widget. A common use-case is to guide new users of a GUI. Many toolkits support the display of interactive hypertext in a tooltip, which allows the user to request additional details.

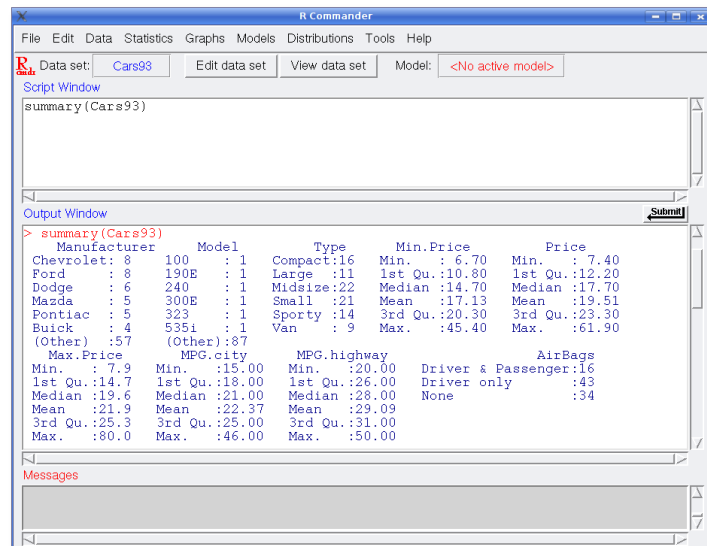


Figure 0.8: Screenshot showing the main Rcmdr (1.3-11) window illustrating the use of multi-line text entry areas for a command area, an output area and a message area.

Progress bars

A progress bar indicates progress on a particular task, which may or may not be bounded. A bounded progress bar usually reports progress in terms of percentage completed. Progress bars should be familiar, as they are often displayed during software installation and while downloading a file. For long-running statistical procedures they can give useful feedback to the user that something is happening.

0.4 Containers

The KDE print dialog of Figure 0.2 contains many of the widgets we discussed in the previous section. Before we can create such a dialog, we need to introduce how to position widgets on the screen. This process is called *widget layout*.

A layout emerges from the organization of the widgets into a hierarchy, where a parent widget positions its children within its allocated space. The top-level window is parentless and forms the root of the hierarchy. A parent visually contains its children and thus is usually called a *container*. This design is natural, because almost every GUI has a hierarchical layout.

It is easy to apply a different layout strategy to each region of a GUI, and when a parent is added or removed from the GUI, so are its children.

It is sometimes tempting for novices to simply assign a fixed position and dimensions for every widget in a GUI. However, such static layouts do not scale well to changes in the state of the application or simply changes to the window size dictated by the window manager. Thus, it is strongly encouraged to delegate the responsibility of layout to a *layout manager* that dynamically calculates the layout as constraints change. Depending on the toolkit, the layout manager might be the container itself, or it might be a separate object to which the container delegates.

Regardless, the type of layout is generally orthogonal to the type of container. For example, a container might draw a border around its children, and this would be independent of how its children are laid out. The rest of this section is divided into two parts: container widgets and layout algorithms. We will continually refer back to the KDE print dialog example as we proceed.

Containers

Top level windows

The top-level window of a GUI is the root of the container hierarchy. All other widgets are contained within it. The conventional main application window will consist of a menubar, a tool bar and a status bar. The primary content of the window is inserted between the tool bar and the status bar, in an area known as the *client area* or *content area*. In the case of a dialog, the content usually appears above a row of buttons, each of which represent a possible response. The print dialog conforms to the dialog convention. The print options fill the content area, and there is a row of buttons at the bottom for issuing a response, such as “Print”.

A window is typically decorated with a title and buttons to iconify, maximize, or close. In the case of the print dialog, the top-level window is entitled “Print – KPDF.”. Besides the text of the title, the decorations are generally the domain of the window manager (often part of the operating system). The application controls the contents of the window.

Once a window is shown, its dimensions are managed by the user, through the window manager. Thus, the programmer must size the window before it becomes visible. This is often referred to as the “default” size of the window. Positioning of a top-level window is generally left to the window manager.

The top-level window forwards window manager events to the application. For example, an application might listen to the window close event in order to prompt a user if there are any unsaved changes to a document.

Tabbed notebooks

A notebook widget depicts each child as if it were a page in a notebook. A page is selected by clicking on a button that appears as a tab. Only a single child is shown at once. The tabbed notebook is a space efficient, categorizing container that is most appropriate when a user is only interested in one page at a time. Modern web browsers take advantage of it to allow several web pages to be open at once within the same window. In the KDE print dialog, detailed options are collapsed into a notebook in order to save space and organize the multitude of options into simple categories: “Copies”, “Advanced Options”, and “Additional Tags”.

Frames

A frame is a simple container that draws a border, possibly with a label, around its child. The purpose of a frame is to enhance comprehension of a GUI by visually distinguishing one group of components from the others. The displayed page of the notebook in Figure 0.2 contains two frames, visually grouping widgets by their function: either Page Selection or Output Settings.

Expanding boxes

An expanding container, or box, will show or hide its children, according to the state of a toggle button. By way of analogy, radio buttons are to notebooks as check buttons are to expanding containers. An expanding box allows the user to adapt a GUI to a particular use case or mode of operation. Often, an expanding box contains so-called “advanced” widgets that are only occasionally useful and are only of interest to a small subset of the users. For example, the Options button in Figure 0.2 controls an expanding box that contains the print options, which are usually best left to their defaults.

Paned boxes

Usually, a layout manager allocates screen space to widgets, but sometimes the user needs to adapt the allocation, according to a present need. For example, the user may wish to increase the size of an image to see the fine details. The *paned container* supports this by juxtaposing panes, either vertically (stacked) or horizontally. The area separating the panes, sometimes called a *sash*, can be adjusted by the user with the mouse.

Layout algorithms

Box layout

The box layout is the most common type of layout algorithm for positioning child components. A box will pack its children either horizontally or vertically¹. Usually, the widgets are packed from left to right, for horizontal boxes, or from top to bottom, in the case of a vertical box. The upper left figure in Figure 0.9 illustrates these possibilities.

The box layout needs to allocate space to its children in both the vertical and horizontal directions. The typical box layout algorithm begins by satisfying the minimum size requirements of its children. The box may need to request more space for itself in order to meet the requirements.

Once the minimum requirements are satisfied, it is conventional and usually desirable for the widgets to fill the space in the direction orthogonal to the packing. For example, widgets in a horizontal box will fill all of their vertical space (the upper right graphic in Figure 0.9 shows some fill possibilities). When this is not desired, most box widgets support different ways of vertically (or horizontally) aligning the widgets (the lower left graphic in Figure 0.9).

More complex logic is involved in the allocation of space in the direction of packing. Any available space after meeting minimum requirements needs to be either allocated to the children or left empty. This depends on whether any children are set to expand. The available space will be distributed evenly to all expanding children. Each child may fill that space or leave it empty. The non-expanding children are simply packed against their side of the container. If there are no expanding children, the remaining space is left empty in the middle (or end if there are no widgets packed against the other side). See the lower right panel in Figure 0.9. One could think of this space being occupied by an invisible spring. Invisible expanding widgets also act as springs.

The button box in the KDE print dialog shows five buttons as child components. At first glance the sizing appears to show that each button is drawn to fully show its label with some fixed space placed between the buttons. If the dialog is expanded, it is seen that there is a spring between the 3rd and 4th buttons, so that the first 3 are aligned with the left side of the window and the last two the right side.

Grid layout

The box layout algorithm only aligns its children along a single dimension. The horizontal box, for example, vertically aligns its children. Nevertheless, nesting permits the construction of complex layouts using only

¹The pack command of `tcltk` can mix the two directions




Figure 0.9: Different possibilities for packing child components within a box. The upper left shows horizontal and vertical layout. The upper right shows some possible alignments or anchorings. The lower left shows that a child could “expand” to fill the space either horizontally, vertically, or both. The lower right shows both a fixed amount of space between the children and an expanding spring between the child components.

simple boxes. However, it is sometimes desirable to align widgets in both dimensions, i.e., to lay them out on a grid. The most flexible grid layout algorithms allow non-regular sizing of rows and columns, as well as the ability for a widget to span multiple cells. Usually, a widget fills the cells allocated to it, but if this is not possible, it may be anchored at a specific point within its cell.

The widgets in the “Printer” frame of Figure 0.9 are subject to a grid layout with five columns and six rows. The first row begins with the “Name:” label, and each widget in that row occupy a separate column. This exposes the size of each column. The first column has only labels, with text justified to the left. The labels are aligned horizontally to each other and vertically with the adjacent field.